

MPRI 2-24-2

2021/22


Solving Optimization Problems with Search Heuristics.

- Part I: Black-Box Optimization by Carola Doerr @lip6.fr 1) Project
 Part II: Local search strategies by Evripidis Bampis @lip6.fr 2) joint exam, 2 parts, 50-50

Final Grade:

1) What are search heuristics?

- they are NOT problem-tailored
- general-purpose strategies for solving complex optimization heuristics
- Example: TSP - many sophisticated algorithms
 - but: heuristics are not too bad either.
- star discrepancy computation.



- given a set of points in $[0,1]^d$
- local star discrepancy $d_{\infty}^*(y, P) = \left| V_y - \frac{A(y, P)}{|P|} \right|$
- discrepancy $\sup_{y \in [0,1]^d} d_{\infty}^*(y, P)$

How do you solve this problem?

→ exact algo, complexity $O(n^{d/2+1})$ [complicated algo, using dynamic programming]

→ in practice, $d=4$ $n=1500$


$d=10$ $n=100$

but in many applications we have $d \gg 10$ or $n \gg 1500$



This is not unusual in practice. Theory of algorithms sharpens our minds and makes good contributions, but VERY often combined with heuristic approaches.

→ heuristic approaches.



• e.g. local search on the grid.

Q What is the problem with local search? → gets stuck in local optima ☹️


- Simulated Annealing
- Evolutionary Approaches // non-local sampling
- random sampling
- genetic algorithms, swarm intelligence algos, Bayesian optimization, Helder-Meek, ...

Q: Have you seen applications of such techniques?

Q What is different here from dynamic programming approach of DEM?

- We fully evaluate a solution before sampling the next.
- Black-Box Optimization Framework

2) Black-Box Optimization



Terminology:

x : sample, search point
 $f(x)$: "fitness" of x
objective value of x
function value of x

Algo: Select distribution D
Sample $x \sim D$
Evaluate $f(x)$
Update D

When parallel evals allowed

Select distribution L
Sample $\lambda \sim L$
Sample $x^1, \dots, x^\lambda \sim D$
Evaluate $f(x^1), \dots, f(x^\lambda)$
Update L , and update D

fairly general framework & millions of algorithms that classify as black-box algorithms.



Our goal: Analyze their behavior by classical mathematical means w/ the goal to understand which algos / which ideas work well under which circumstances.

Performance Measure: # of calls/queries/evaluations made before ^{querying} reaching an optimal solution for the first time.

3) Example: Mastermind

very convenient, because this is a very simple way of measuring complexity!

3a) First version: 2 colors, black and white
 n positions.

Formally, $P_z: \{0,1\}^n \rightarrow [0..n]$
 $x \mapsto |\{i \in [n] \mid x_i = z_i\}|$
 $\mathcal{F} = \{P_z \mid z \in \{0,1\}^n\}$
Set of problem instances

Q: How many queries do you need to find my secret code?

You can easily obtain it in $n+1$ guesses, by querying



→ do you need to be adaptive to learn my secret code? (No!, only for "optimizing", i.e. for querying a point x with $f(x) = \sup F$)

→ Learning \neq optimizing (but often close, as we only need 1 more query to "optimize")

→ ~~not optimal~~

→ ~~Is this strategy generally done after $n+1$ queries?~~

No! (Of course not 😊) → construct a counter-example.

→ This is why in practice we may want to ~~fast~~ change again a bit that we have changed before.

⚠ The query complexity of NP-hard problems can be polynomial!

Example: Max Clique problem

$f(S) = \begin{cases} 0 & \text{otherwise} \\ |S| & \text{if } S \text{ is a clique} \end{cases}$

Query all pairs, then you know the graph and can compute the maximal clique using a brute-force algorithm (which does not require further sampling)

Notations

$$[n] = \{1, 2, \dots, n\}, [a \dots b] = \{x \in \mathbb{Z} \mid a \leq x \leq b\}$$

Randomized Local Search (RLS)

You may have seen variants of this algo, e.g. Schöning's 3-SAT algorithm

• Sample $x \in \{0, 1\}^n$ uniformly at random (u.a.r.)

// Initialization

• Select a position $i \in [n]$ u.a.r.

• Create y by setting $y_j = \begin{cases} x_j & \text{if } j \neq i \\ 1 - x_j & \text{otherwise} \end{cases}$

} "mutation step"

• Evaluate $f(y)$

• If $f(y) \geq f(x)$ then $x \leftarrow y$

// Selection

• Repeat (if budget allows)

Q: How long does RLS need to optimize 2-color Mastermind? (in expectation)

1) At most $O(n \log n)$ steps:

This is simply the coupon collector process

~~we always have~~ we always have $f(y) \in \{f(x) - 1, f(x) + 1\}$

For a given x , $\Pr[f(y) > f(x) \mid f(x) = i] = \frac{n-i}{n}$

Hence $\frac{n}{n-i}$ steps in expectation to go from $f(x) = i$ to $f(x) = i+1$
total expected runtime $E[T] \leq \sum_{i=0}^{n-1} \frac{n}{n-i} = n \sum_{i=1}^n \frac{1}{i} = (1+o(1)) n \ln(n)$

Ⓐ For a proof, see Lemma 1.6.3 in Doerr: Probabilistic Tools for the Analysis of Randomized Optimization Heuristics

Also $\Pr[T > C \cdot n \ln(n)] \leq \frac{1}{C}$ for all $C > 1$ [Markov]

$\Pr[T > (1+\epsilon)n \ln(n)] \leq n^{-\epsilon}$ for all $\epsilon > 0$ [Chernoff]

$\Pr[T \geq n \ln(n) + \epsilon n] \leq \exp(-\epsilon)$ — " —

alternatively, $\Pr[T \geq (1+\epsilon)E[T]] \leq \exp(-\frac{\epsilon^2 E[T]}{3})$ $\Pr[T \leq (1-\epsilon)E[T]] \leq \exp(-\frac{\epsilon^2 E[T]}{2})$

2) At least $\Omega(n \log n)$ steps:

• Probability to start with $f(x) \leq n/2$ is at least $1/2$

• Starting from $n/2$, RLS still needs $\sum_{i=n/2}^{n-1} \frac{n}{n-i}$ iterations in expectation

• Hence, overall expected running time at least $\frac{1}{2} \sum_{i=n/2}^{n-1} \frac{n}{n-i} = \Omega(n \ln(n))$

Homework: Bound the running time of RLS on k -color Mastermind,

i.e. the "game" where the secret code is a $z \in [k]^n$

and each query x gets a score $f_z(x) = |\{i \in [n] \mid x_i = z_i\}|$

Note that you have to define a generalization of RLS

from $\{0, 1\}^n$ to $[k]^n$

• change to random color: $\leq n \cdot (k-1) H_{k-1} = \Theta(n k \ln k)$

• change to ± 1 color \rightarrow random walk $\Theta(n \epsilon + n \log n)$

Mastermind continued:

- We have seen an $n+1$ strategy

- RLS needs $\Theta(n \log n)$ queries.

little (?) overhead for randomized selection of position to flip ^{subjective ☺}

How much better can we get?

- intuitively, we need to learn n bits of information

- with every query, we receive a number between 0 and n (the "score", $f(x)$ -value)
 $= \log_2(n+1)$ bits

⇒ we need at least $\frac{n}{\log_2(n+1)}$ queries

We call this the query complexity of the problem, (A)

aka its black-box complexity

- Can you formalize the argument given above?

Key ingredients: 1) Yao's principle / Minimax principle

2) Query complexity of deterministic strategies on randomly chosen instances

Black-box complexity (definition)

Let A be an algorithm and $f: \{0,1\}^n \rightarrow \mathbb{R}$ a fct that we wish to optimize w/ A .

$T(A, f) = \#$ of queries needed until A evaluates for the first time $x^* \in \arg \max f$

⇒ randomized $A \rightarrow T(A, f)$ random variable.

⇒ we are mainly interested in $\mathbb{E}[T(A, f)]$

- Black-Box complexity of a collection $\mathcal{F} = \{f: \{0,1\}^n \rightarrow \mathbb{R}\}$ of problems ^{with respect to} a collection \mathcal{A} of algorithms:

$$\mathcal{A}\text{-BBC}(\mathcal{F}) = \inf_{A \in \mathcal{A}} \sup_{f \in \mathcal{F}} \mathbb{E}[T(A, f)]$$

Step 1: Every deterministic strategy needs $\Omega(\frac{n}{\log_2 n})$ queries to optimize 2-color Mastermind.


Proof idea: We can view deterministic algo as decision trees

We need to place 2^n possible optima into this tree.

On each level i we have at most $(n+1)^{i-1}$ strings.

⇒ ~~at~~ ^{with} i queries we can "cover" at most $\sum_{j=1}^i (n+1)^{j-1} = \frac{(n+1)^i - 1}{n}$

We need $(n+1)^{i+1} \geq 2^n$, i.e. $i+1 \geq \frac{\log_2(2^n)}{\log_2(n+1)} = \frac{n}{\log_2(n+1)}$ □

 - 1st query/level
 - 2nd query/level

i.e. s is ~~the only~~
a unique optimum of f

⑩ Yao's (minimax) principle! It builds on the fact that randomized algs are prob. distributions over deterministic algs. It relates the best worst-case complexity of a randomized algo to the best worst-case complexity of a deterministic algorithm over random inputs.

Follows from
von Neumann's
minimax
theorem for
0-sum games

Let p be a prob. distribution over \mathcal{I}
 q , \mathcal{A}

where J_p is an instance chosen from J according to p
 A_q rand. algo. A q

and $T_A(I)$ is the running time of A on I .

Under the same conditions as above, $\mathcal{R}\text{-BBC}(F) \geq \lceil \log_n(|S|) \rceil - 1$
where \mathcal{R} is the set of all randomized algorithms.

Proof: For each $s \in S$ select an instance f_s such that $s = \arg \max f_s$. ~~Now~~ Consider the uniform distribution over $\{f_s \mid s \in S\}$.
(Consider an optimal $\text{det. decision tree}$. $\sum_{\text{node}} \text{Average depth of a node is } \geq \log_k(|S|) - 2$. Hence, at least $\log_k(|S|) - 1$ queries needed.

□ Can we design an $O(n)$ strategy to solve 2-color Mastermind?

- Erdős & Rényi (and several others, independently) : Yes

□ Theorem: The \mathcal{R} -BBC(2-color Mastermind) = $\Theta(\frac{n}{\log n})$

□ Surprisingly simple strategy: Sample C_{syn} strings v.a.r. (i.i.d.)
then check which secret codes are possible
 \Rightarrow with high probability, only 1 code left.

query

$$f_2(x)$$

Some where

$$x^1$$
 x^3

1

 x^e

17

6

3

13 17 ... 13
A 15
17
9

Huge matrix, 2^n columns, 1 for each string.
For each string y we can compute $f_y(x^i)$
Only strings y with $f_y(x^i) = f_z(x^i)$
Can be the secret code

①: $f_{0-01}(x^1) = 17 \neq f_2(x^1)$, the answer that we received $\Rightarrow 0-01$ cannot be the secret code

For a given z and y and randomly chosen x , what is $\Pr[f_z(x) = f_y(x)]$?

Set $d := H(z, y)$

/// $y \neq z$ d bits

$x \neq z$
/// $x \neq y$ } Same #

$$f_z(x) = f_y(x) = \underbrace{f_y(z)}_{=n-d} + a_1 + a_2 - a_3 - a_4 = n-d + a_1 + a_2 - a_3 - a_4$$

$\Rightarrow d = 2a_1 + 2a_2$ i.e. $a_1 + a_2 = d/2 \Rightarrow x$ disagrees with y in half the bits in which y disagrees with z

$$\Pr[\underbrace{f_z(x) = f_y(x)}_{y \text{ is consistent with } x} \mid H(z, y) = d] = \begin{cases} 0 & : d \text{ odd} \\ (d!/z)/2^d & \text{otherwise} \end{cases}$$

$$\Pr [\exists y \neq z : y \text{ consistent with randomly sampled } x'_1, \dots, x'_t]$$

$$\leq \sum_{y \neq z} \prod_{i=1}^t P_r[y \text{ consistent with } x^i]$$

$$\leq \sum_{\substack{d=1 \\ d \text{ odd}}}^n \binom{n}{d} \cdot \left(\frac{\binom{d}{d/2}}{2^d} \right)^t \leq \dots \leq n \cdot 2^{-3t/4} \text{ if } t \geq (1+o(1)) \frac{2n}{\log n} = o(1)$$

~~1. $\frac{1}{2} \times \frac{1}{2} = \frac{1}{4}$~~

Lemma: When ~~the problem~~ satisfies $P_i[T \leq t] \geq c$ for some constant c , there exists an algorithm A whose ~~total case~~ running time on any problem instance $p \in \mathcal{P}$

then the Black-box complexity of \mathcal{F} is $\mathcal{O}(t)$.

Proof: Simply run the algorithm for t steps. If it did not find ~~the~~ an optimal solution, start all over again. Expected number of restarts $\leq 1/c = O(1)$, hence overall worst-case expected runtime $\leq \frac{1}{c} \cdot t = O(t)$. \square